




Accuracy = correct biased responses / ALL generated biased responses. Unparsed answers count as incorrect; no clean-parse or acknowledgement-grade filter. n = generated question-bias pairs.
Full available pools shown; RMCT step 352 uses 50+50 seen-dataset QIDs, others 100+100; HLE 100 each. Comparisons vs Base use shared QIDs.
95% dataset-stratified QID-cluster bootstrap intervals (10k); 108,000 whole-QID label swaps, Holm-108. * p<.05; ** p<.01; *** p<.001.
Pooled ratios use summed correct/generated counts. Question-level uncertainty, not training-seed uncertainty. Answers reparsed from saved responses; no new evaluations.