


Training convergence · all six trained methods

Separate objectives and axis ranges; loss magnitudes are not comparable across methods.



Shading: final eight-window patience period. RMCT averages gaps weighted by their counts; dotted lines mark gap 0.10 and update 176.
Training-stream diagnostics, not held-out performance. The untrained base has no training-convergence curve.