



AITA-NTA-FLIP · 64k retries · terminal comparison


Observed both-NTA / all 1,591 pairs — invalid pairs cannot count as both-NTA



Both-NTA / 1,194 pairs parsed by EVERY displayed condition — selected subset



Remaining invalid responses: all reached 65,536 tokens



Thinking enabled. Only 20k cap hits retried at 64k; other responses unchanged.

Step numbers under methods. 95% paired-bootstrap CI (10,000 resamples).

Stars vs Base: exact McNemar + Holm within each panel (* < .05, ** < .01, *** < .001).

Lower both-NTA suggests less sycophancy, but missing verdicts confound the all-pairs rate; complete-case selection is not a correction.