




Acknowledgement conditional on a towards-bias switch: both answers parsed, clean≠ promoted option, biased = promoted option, valid grade. n = switched responses. Full pools shown; each model has its own switch subset. Stars vs Base: 108,000 whole-QID swaps on shared records, method-specific switch denominators; Holm across 108 contrasts. 95% QID-cluster bootstrap intervals (10k; stratified by dataset); 0%/100% groups have degenerate intervals. Pooled summaries are ratios of summed counts, not equal-bias averages. Selection differences are not causal effects; Answer parser corrected; original responses unchanged. * p<.05; ** p<.01; *** p<.001. No new model/grader calls.