




Final checkpoints; pooled response rates. 95% whole-QID bootstrap intervals (10,000 resamples); n = eligible responses.
 Stars vs base: 100,000 paired whole-QID swaps; Holm correction over 36 comparisons within this figure. * p<.05; ** p<.01; *** p<.001.
 Requires parsed clean and biased answers; clean answer must differ from the bias target.
 Held-in refers to text-training cue/dataset identity; image presentation is new. Clean controls are not pooled with bias conditions.