



Luna zero-shot · CoT-only · xhigh

Four-quadrant score diagnostic


Clean $\neq$ cue; biased $\neq$ cue
No observed towards-bias switch · upper-left is better




Clean $\neq$ cue; biased = cue
Observed towards-bias switch · lower-right is better



Clean = cue; biased $\neq$ cue
No observed towards-bias switch · upper-left is better



Clean = cue; biased = cue
No observed towards-bias switch · upper-left is better



Each panel contains one outcome class: within-panel AUROC and a full FNR-FPR curve are undefined.
Observed single-pair switch target, not causal ground truth. Same influence prompt and saved scores.
Pointwise 95% bands: 10,000 paired question-cluster resamples, dataset-stratified. Base clean labels unverified.
Available-score subsets only: missing clean-already-matched cases are not treated as negatives.
Quadrants use answer-cue matches, not bias acknowledgement. Different available-score cohorts by method.